

Clinical Statistics Methods Forum Data Challenge

Eric Polley & Prabin Thapa

September 20th, 2016

Introduction

The 2016 Clinical Statistics Methods Forum Data Challenge

The goal is to construct a predictor of therapeutic drug dose given baseline clinical measurements. Participants given a training data set with 3,000 patients and will be evaluated on a blinded validation data set with 1,722 patients

Overview

- ▶ Working in teams, $n \in (1, 2, \dots, 20)$, use the training data to construct a predictor for the therapeutic dose given baseline clinical measurements
- ▶ Predictors evaluated by mean squared error of predicted dose and true dose
- ▶ With validation data set, provide predictions for each individual. A file with the Subject ID and the predicted dose as 2 columns can be email to me
- ▶ Don't forget to include your team name
- ▶ On Oct. 18th, meet for a midpoint review and group discussion
- ▶ On Nov. 15th, each team will be asked to provide a short summary on how they constructed the predictor and I will reveal the performance on the validation data
- ▶ Everyone is welcome to attend the discussion (and drink coffee)
- ▶ Prizes to be determined

Overview


- ▶ 2 data sets provided in comma separated format
- ▶ TRAIN.CSV includes the baseline variables and outcome of interest (Therapeutic_Dose)
- ▶ VALID.CSV includes the baseline variables, need to predict the outcome
- ▶ Data available on GitHub:
https://github.com/ecpolley/CSMF_Data_Challenge

How To Get Data

```
# link to data on GitHub page if not available
if(file.exists("TRAIN.CSV")) {
  TRAIN <- read.csv("TRAIN.CSV")
} else {
  urlfile <- "https://raw.githubusercontent.com/ecpolley/
    CSMF_Data_Challenge/master/TRAIN.CSV"
  download.file(urlfile, destfile = "TRAIN.CSV")
  TRAIN <- read.csv("TRAIN.CSV")
}
dim(TRAIN)
```

```
## [1] 3000 44
```

Therapeutic_Dose





Baseline Variables

```
str(TRAIN[, 1:8])
```

```
## 'data.frame':    3000 obs. of  8 variables:
## $ Subject_ID : Factor w/ 3000 levels "W0126718362",...: 1013 2027 7
## $ Project.Site: int  4 3 7 1 6 7 16 5 7 3 ...
## $ Gender      : Factor w/ 2 levels "female","male": 1 2 1 2 2 1 1 2
## $ Race        : Factor w/ 4 levels "Asian","Black or African Americ
## $ Ethnicity   : Factor w/ 3 levels "Hispanic or Latino",...: 2 2 3 2
## $ Age         : Factor w/ 9 levels "10_19","20_29",...: 5 5 6 7 7 8
## $ Height      : num  155 166 144 170 158 ...
## $ Weight      : num  59 70 55.3 118.4 84 ...
```

Baseline Variables



Baseline Variables

```
str(TRAIN[, 9:12])
```

```
## 'data.frame':    3000 obs. of  4 variables:
## $ Indication_for_Treatment      : Factor w/ 41 levels
## $ Diabetes                      : int   1 NA 0 NA NA
## $ Congestive_Heart_Failure_and_or_Cardiomyopathy: int   0 0 0 NA NA
## $ Valve_Replacement             : int   0 1 0 NA NA
```

Baseline Variables

```
str(TRAIN[, 13:29])
```

```
## 'data.frame':    3000 obs. of  17 variables:
```

```
## $ Aspirin                : int  0 NA 0 0 NA 0 1 0 0
## $ Acetaminophen_or_Paracetamol : int  0 NA 1 NA NA NA NA
## $ Simvastatin             : int  0 NA 0 0 NA 0 0 0 0
## $ Atorvastatin            : int  0 NA NA NA NA NA 1
## $ Fluvastatin             : int  0 NA NA NA NA NA NA
## $ Lovastatin              : int  0 NA NA NA NA NA NA
## $ Pravastatin             : int  0 NA NA NA NA NA NA
## $ Rosuvastatin            : int  0 NA NA NA NA NA NA
## $ Cerivastatin            : int  0 NA NA NA NA NA NA
## $ Amiodarone              : int  0 NA 0 0 NA 0 1 0 0
## $ Carbamazepine           : int  0 NA NA NA NA NA NA
## $ Phenytoin               : int  0 NA NA NA NA NA NA
## $ Rifampin                : int  0 NA NA NA NA NA NA
## $ Sulfonamide_Antibiotics : int  0 NA NA NA NA NA NA
## $ Macrolide_Antibiotics   : int  0 NA NA NA NA NA NA
## $ Anti_fungal_Azoles      : int  0 NA NA NA NA NA NA
## $ Herbal_Medications_Vitamins_Supplements: int  1 NA NA NA NA NA NA
```

Baseline Variables

```
str(TRAIN[, 30:33])
```

```
## 'data.frame':    3000 obs. of  4 variables:
##  $ Current_Smoker : int  NA NA NA 0 NA 0 1 0 NA NA ...
##  $ GeneA           : Factor w/ 11 levels "*1/*1","*1/*11",...: 1 1 1 1
##  $ GeneB           : Factor w/ 3 levels "A/A","A/G","G/G": 2 1 2 2 1
##  $ NoComorbidities: int   0 0 0 0 1 0 0 0 0 0 ...
```

Baseline Variables

```
str(TRAIN[, 34:43])
```

```
## 'data.frame':    3000 obs. of  10 variables:
## $ Biomarker_1 : num  -4.76 -9.35 -15.28 -9.45 -10.71 ...
## $ Biomarker_2 : num  -1.608 3.454 5.451 -0.612 2.161 ...
## $ Biomarker_3 : num  1.477 0.463 2.413 2.913 1.268 ...
## $ Biomarker_4 : num  4.55 3.42 4.18 4.85 5.56 ...
## $ Biomarker_5 : num  2.41 -1.61 3.74 4.42 2.18 ...
## $ Biomarker_6 : num  3.3 -1.15 6.26 5.39 3.03 ...
## $ Biomarker_7 : num  1.885 1.755 1.731 1.582 0.114 ...
## $ Biomarker_8 : num  -7.21 -7.01 -4.2 -5.01 -5.88 ...
## $ Biomarker_9 : num  -0.1227 0.2597 -0.0925 -0.1557 -1.0295 ...
## $ Biomarker_10: num  5.22 8.36 8.26 7.31 7.17 ...
```

Basic Regression Predictor

```
fit <- lm(Therapeutic_Dose ~ Age + Gender +  
          Biomarker_1 + Biomarker_2, data = TRAIN)  
fit
```

```
##
```

```
## Call:
```

```
## lm(formula = Therapeutic_Dose ~ Age + Gender + Biomarker_1 +  
##     Biomarker_2, data = TRAIN)
```

```
##
```

```
## Coefficients:
```

```
## (Intercept)      Age20_29      Age30_39      Age40_49      Age50_59  
##    42.09257     -2.81562     -2.43856     -3.97531     -6.98273  
##   Age60_69     Age70_79     Age80_89      Age90+  Gendermale  
##  -12.66090    -16.90053    -16.59322    -21.00917      2.24424  
## Biomarker_1  Biomarker_2  
##     0.03230     -0.08105
```

How To Get Data

```
# link to data on GitHub page if not available
if(file.exists("VALID.CSV")) {
  VALID <- read.csv("VALID.CSV")
} else {
  urlfile <- "https://raw.githubusercontent.com/ecpolley/
    CSMF_Data_Challenge/master/VALID.CSV"
  download.file(urlfile, destfile = "VALID.CSV")
  VALID <- read.csv("VALID.CSV")
}
dim(VALID)
```

```
## [1] 1722 43
```

Basic Regression Predictor

```
pred_lm <- predict(fit, newdata = VALID)
pred_lm_df <- data.frame(ID = VALID$Subject_ID, predict = pred_lm)
head(pred_lm_df)
```

```
##           ID predict
## 1 W0151579099 31.46625
## 2 W0151957677 37.27828
## 3 W0151693136 26.53852
## 4 W0135312561 28.54693
## 5 W0151752657 24.78085
## 6 W0150479653 41.08882
```

```
# write out prediction table
# write.csv(pred_lm_df, "final_predictions.csv")
```

Cross-Validation

Prior to submitting predictions, should evaluate performance

```
## get V-fold CV estimate of MSE
V <- 10
N <- nrow(TRAIN)
MSE_cv <- rep(NA, V) # placeholder for CV MSE estimates
# list of row ids by V validation splits
validRows <- split(sample(1:N), rep(1:V, length=N))
for(v in seq(V)) {
  tempTRAIN <- TRAIN[-validRows[[v]], ]
  tempVALID <- TRAIN[validRows[[v]], ]
  fit_cv <- lm(Therapeutic_Dose ~ Age + Gender +
               Biomarker_1 + Biomarker_2, data = tempTRAIN)
  pred_cv <- predict(fit_cv, newdata = tempVALID)
  MSE_cv[v] <- mean((pred_cv - tempVALID$Therapeutic_Dose)^2)
}
```


Cross-Validation

```
# get V-fold CV estimate of MSE
```

```
summary(MSE_cv)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    215.9   223.4   238.3   243.1   260.0   278.1
```

```
# CV R-squared
```

```
1 - mean(MSE_cv)/var(TRAIN$Therapeutic_Dose)
```

```
## [1] 0.0933959
```

Questions?